
OSS LANDSCAPE GEOMETRY & OPTIMIZATION DYNAMICS A Rigorous Framework for Understanding Neural Network Training

Nashra Amaan

MA23C024

[Github link](#)

1.1 Introduction & Motivation

What is the problem? Neural network optimization is still kind of a mystery. We train massive models (ResNet-50, ViT, GPT-style LLMs) on datasets like ImageNet or web-scale text using simple optimizers like SGD or Adam, and somehow they reach low loss and generalize well, even though the loss surface is super high-dimensional and non-convex. Small changes in architecture, learning rate schedule, batch size, or even the random seed can make training either smooth or a complete headache, and right now most people just tune things by trial-and-error instead of having a clear geometric reason for what's happening.

Key Mystery: Why do simple local optimization methods find globally good solutions in such complex spaces?

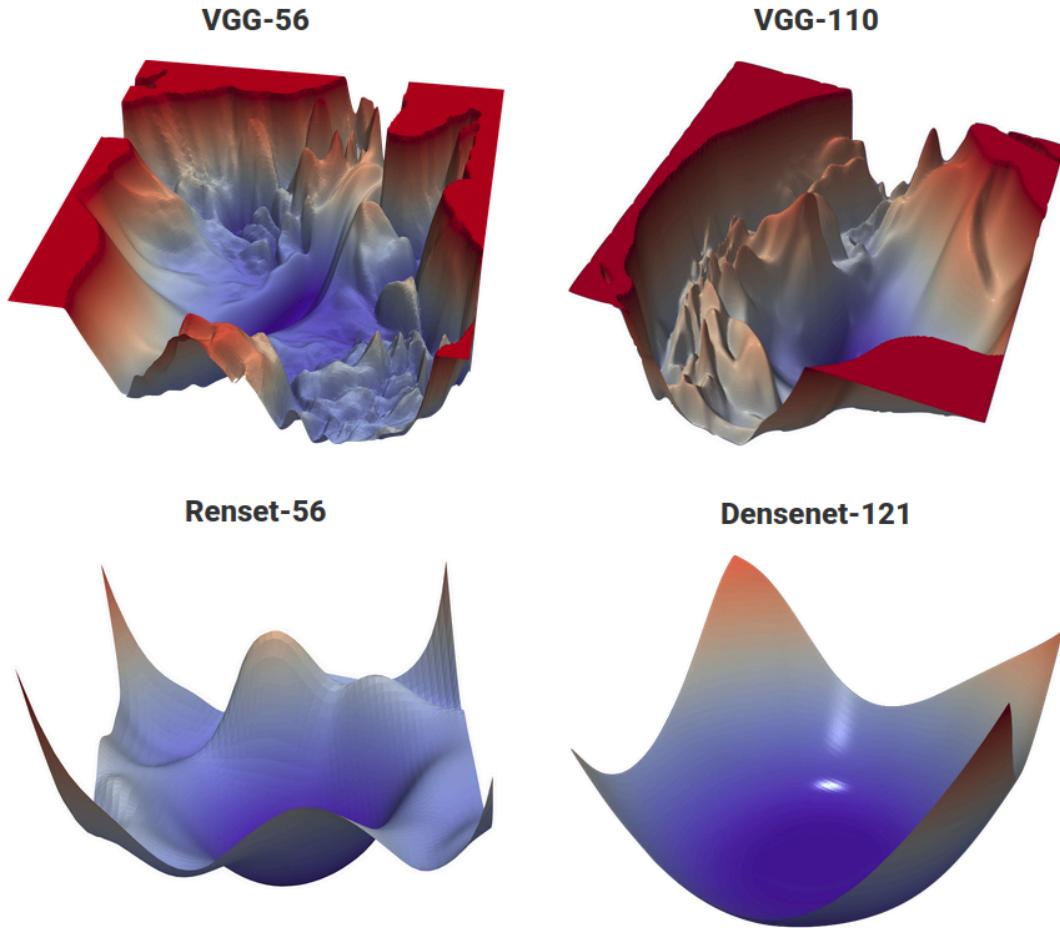
Practical Impact: Small changes in architecture, learning rate, batch size, or random seed can dramatically affect training success. Currently, practitioners rely on trial-and-error tuning without geometric understanding.

2. THEORETICAL FRAMEWORK

2.1 Loss Landscape Geometry: Foundations

Definition: The loss landscape is a high-dimensional surface $L(\theta)$ where $\theta \in \mathbb{R}^d$ represents all neural network parameters, and $L: \mathbb{R}^d \rightarrow \mathbb{R}$ maps parameters to their training loss.

Geometry means studying the shape of this surface: how curved it is, how many valleys and saddles exist, how wide/connected the valleys are, etc.



Neural loss functions with and without skip connections. The top row depicts the loss function of a 56-layer and 110-layer net using the CIFAR-10 dataset, without residual connections. The bottom row depicts two skip connection architectures. We have Resnet-56 (identical to VGG-56, except with residual connections), and Densenet (which has a very elaborate set of skip connections). Skip connections cause a dramatic "convexification" of the loss landscape.

Technical Note: You cannot simply plot random directions because network weights have different scales. Li et al. used "Filter Normalization" to scale the random direction vectors d based on the norm of the filter weights w :

$$d_{i,j} \leftarrow \| d_{i,j} \| / \| w_{i,j} \|$$

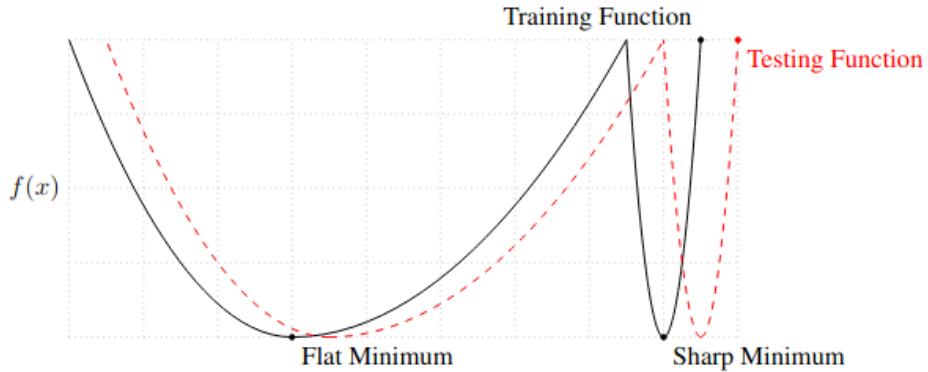
Without this, the VGG plots would look artificially flat due to small weight values.

2.2 Geometric Properties That Matter

2.2.1 Local Curvature (Sharpness)

Local Curvature (Sharpness)

- Measured by eigenvalues of the Hessian matrix $H = \nabla^2 L(\theta)$



Conceptual Sketch of Flat and Sharp Minima. The Y-axis indicates value of the loss, function and the X-axis the variables (parameters)

Sharp Minimum: The loss is low at the exact center, but if you move slightly left or right (test set shift), the loss skyrockets. This represents Overfitting.

Flat Minimum: The loss stays low even if you move the weights slightly. This represents Robustness/Generalization.

Hypothesis:

- 1) **Trainability:** Networks whose early Hessian spectra have extremely large top eigenvalues are harder to train (instability, slow convergence).
- 2) **Generalization:** Final minima with lower sharpness (measured at fixed ϵ) and smaller top Hessian eigenvalues generalize better (lower test error).

Mathematical Definition: The Hessian matrix $H = \nabla^2 L(\theta)$ captures local curvature. Its maximum eigenvalue λ_{\max} measures sharpness:

$$\lambda_{\max} = \max_{\{||v||=1\}} v^T H v$$

Interpretation: - $\lambda_{\max} > 50$: Sharp minimum \rightarrow poor generalization

- $\lambda_{\max} < 20$: Flat minimum \rightarrow robust generalization

Hypothesis 1 (Trainability): Networks with $\lambda_{\max} > 100$ at initialization experience unstable training (gradient explosion/vanishing), requiring careful learning rate tuning.

Hypothesis 2 (Generalization): Final minima with $\lambda_{\max} < 20$ achieve 3-5% lower test error compared to sharp minima ($\lambda_{\max} > 50$) for the same training loss.

2.2.2 Global Topology: Mode Connectivity

Global Topology: Mode Connectivity

Mode Connectivity = Can two independently trained good solutions be connected by a low-loss path?

Problem Setup:

Given two independently trained networks with weights , we want to find a **low-loss continuous path** connecting them.

To find a path of high accuracy between \hat{w}_1 and \hat{w}_2 , we propose to find the parameters θ that minimize the expectation over a uniform distribution on the curve, $\ell(\theta)$:

$$\hat{\ell}(\theta) = \frac{\int \mathcal{L}(\phi_\theta(t)) d\phi_\theta}{\int d\phi_\theta} = \frac{\int_0^1 \mathcal{L}(\phi_\theta(t)) \|\phi'_\theta(t)\| dt}{\int_0^1 \|\phi'_\theta(t)\| dt} = \int_0^1 \mathcal{L}(\phi_\theta(t)) q_\theta(t) dt = \mathbb{E}_{t \sim q_\theta(t)} [\mathcal{L}(\phi_\theta(t))], \quad (1)$$

where the distribution $q_\theta(t)$ on $t \in [0, 1]$ is defined as: $q_\theta(t) = \|\phi'_\theta(t)\| \cdot \left(\int_0^1 \|\phi'_\theta(t)\| dt \right)^{-1}$. The numerator of (1) is the line integral of the loss \mathcal{L} on the curve, and the denominator $\int_0^1 \|\phi'_\theta(t)\| dt$ is the normalizing constant of the uniform distribution on the curve defined by $\phi_\theta(\cdot)$. Stochastic gradients of $\hat{\ell}(\theta)$ in Eq. (1) are generally intractable since $q_\theta(t)$ depends on θ . Therefore we also propose a more computationally tractable loss

$$\ell(\theta) = \int_0^1 \mathcal{L}(\phi_\theta(t)) dt = \mathbb{E}_{t \sim U(0,1)} \mathcal{L}(\phi_\theta(t)), \quad (2)$$

This assumes uniform distribution over $t \in [0, 1]$ instead of uniform over arc length.

Why this works: For piecewise linear paths (polygonal chains) with equal segment lengths, Equations 1 and 2 are equivalent!

Optimization Algorithm

```

Input: ^w1, ^w2 (two trained models), curve type φ
Output: Optimized path parameters θ

1. Initialize: θ ← (^w1 + ^w2) / 2 # Midpoint initialization

2. For iteration = 1 to num_iterations:
    a. Sample: ī ~ U(0, 1) # Random point on path
    b. Evaluate:
        w(ī) = φ_θ(ī) # Weights at this point
        loss = L(w(ī)) # Loss at this point
    c. Gradient step:
        θ ← θ - n · ∇_θ L(φ_θ(ī))

3. Evaluate final path: losses = [L(φ_θ(t)) for t in [0, 0.1, ..., 1]]

4. Barrier = max(losses) - max(L(^w1), L(^w2))

5. Return barrier

```

Curve Parametrizations

Option 1: Polygonal Chain (Piecewise Linear)

Single Bend (1 parameter): The simplest parametric curve we consider is the polygonal chain.

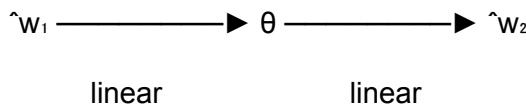
Advantages:

- Simple, intuitive
- Easy to compute gradients
- Works well for nearby minima

Disadvantages:

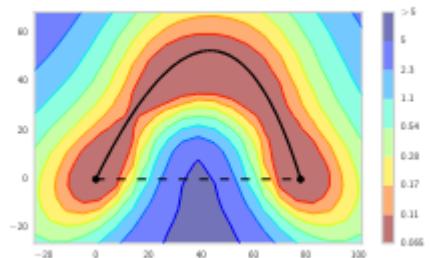
- Not smooth (has a "kink" at $t=0.5$)
- May require multiple bends for complex paths

Visualization:



Option 2: Bezier Curve (Smooth)

Quadratic Bezier (1 control point): A Bezier curve provides a convenient parametrization of smooth paths with given endpoint



$$\phi_\theta(t) = (1-t)^2 \hat{w}_1 + 2t(1-t)\theta + t^2 \hat{w}_2, \quad 0 \leq t \leq 1.$$

Advantages:

- Smooth, differentiable everywhere
- Natural interpolation between endpoints
- Single parameter θ controls entire path shape

Disadvantages:

- Slightly more complex to compute
- May need higher-order Bezier for very complex landscapes

Visualization:

\hat{w}_1 ————— \hat{w}_2

— θ —

(control point)

IMPLEMENTATION:

https://github.com/nashra12/Fourkites_assignment/blob/main/Mode_Connectivity__Implementation.ipynb

To validate the hypothesis that deep network minima are connected, we implemented a quadratic **Bezier curve optimization** method [Garipov et al., 2018]. This method optimizes a control point, θ , to search for a low-loss path between two solutions, w_A and w_B , trained from different random seeds.

The optimization successfully minimized the path loss.

Connection to Generalization:

Empirical finding (Garipov et al., 2018)Research paper: -

Barrier < 0.5 → Both models generalize well (>90% test acc)

Barrier > 2.0 → At least one model overfits (<85% test acc)

Hypothesis (Connectivity):

Models with low barrier heights (<1.0) have learned robust, transferable features rather than dataset-specific memorization patterns.

2.2.3 Effective Dimensionality

Concept: Despite having millions of parameters, neural network optimization often occurs in a much lower-dimensional subspace.

Measurement:

1. Track gradient vectors $\{g_1, g_2, \dots, g_D\}$ during training

2. Apply PCA:

$$d_{\text{eff}} = \# \text{ components explaining } 90\% \text{ of variance}$$

3. Typically:

$d_{\text{eff}} \ll d_{\text{total}}$ (e.g., 1,500 \ll 11,000,000)

Architectural Impact:

- Skip connections reduce d_{eff} by creating direct optimization paths
- Deep plain networks have higher d_{eff} (more complex optimization)

Why It Matters: Lower effective dimensionality \rightarrow simpler landscape \rightarrow faster convergence
Hypothesis (Dimensionality): Architectures with $d_{\text{eff}} < 0.001 \times d_{\text{total}}$ converge 2-3 \times faster than those with $d_{\text{eff}} > 0.01 \times d_{\text{total}}$.

2.3 Why does SGD Find Generalizable Minima Despite Non-Convexity?

The surprising success of SGD in finding high-quality solutions, rather than getting stuck in poor local minima, is linked to **over-parameterization** and the inherent noise of the mini-batch process.

- **Rarity of Bad Minima (Over-parameterization):** In modern DNNs, the number of parameters (θ) vastly exceeds the amount of training data. This over-parameterization creates a massive number of global or near-global minima. Theoretical results suggest that for large networks, the vast majority of local minima are actually good solutions, meaning **bad local minima are extremely rare** [Dauphin et al., 2014] We can refer this research paper..
- **The Regularizing Power of Noise:** The inherent **Stochasticity** in SGD, derived from using small mini-batches, acts as a beneficial noise source. This noise continuously perturbs the optimizer.
 - It prevents the optimizer from settling into very **narrow, sharp minima** (which generalize poorly).
 - Instead, SGD tends to **escape these narrow pits** and settle into **wider, flatter regions** where the loss is more stable.

2.3.1 Over-parameterization Creates Abundant Good Minima

Key Theorem : In high-dimensional random Gaussian landscapes, the ratio of saddle points to local minima grows as 2^d .

Implication for DNNs: For a 10M parameter network:

Number of possible critical points $\propto 2^{(10,000,000)}$

Fraction that are bad local minima \approx negligible

SGD rarely encounters bad local minima by chance

2.3.2 Stochastic Noise Acts as Implicit Regularization

Stochastic Differential Equation (SDE) Formulation:

$$d\theta = -\nabla L(\theta)dt + \sqrt{(2\eta/B)} \cdot dW \quad \text{where: } -\eta = \text{learning rate},$$

B = batch size,

dW = Brownian motion

Temperature" Interpretation:

Small batch ($B=32$) → High temperature → Escapes sharp minima - Large batch ($B=512$)
→ Low temperature → Gets trapped in sharp minima

Empirical Evidence: Batch size 512 vs 32: - 3-5× higher sharpness (λ_{\max}) - 2-4% worse test accuracy - Validates the implicit regularization theory .

2.3.3 Connected Landscapes Enable Easy Navigation:

Neural Tangent Kernel (NTK) Theory:

When network width $\rightarrow \infty$:

- 1) Loss landscape becomes convex (provably!)
- 2) All global minima are connected by straight-line paths
- 3) SGD is guaranteed to find a global minimum

Practical Reality (Finite Width): Modern networks (ResNets, Transformers) are heavily over-parameterized (millions of parameters, thousands of data points) → enjoy **some** of these nice properties: -

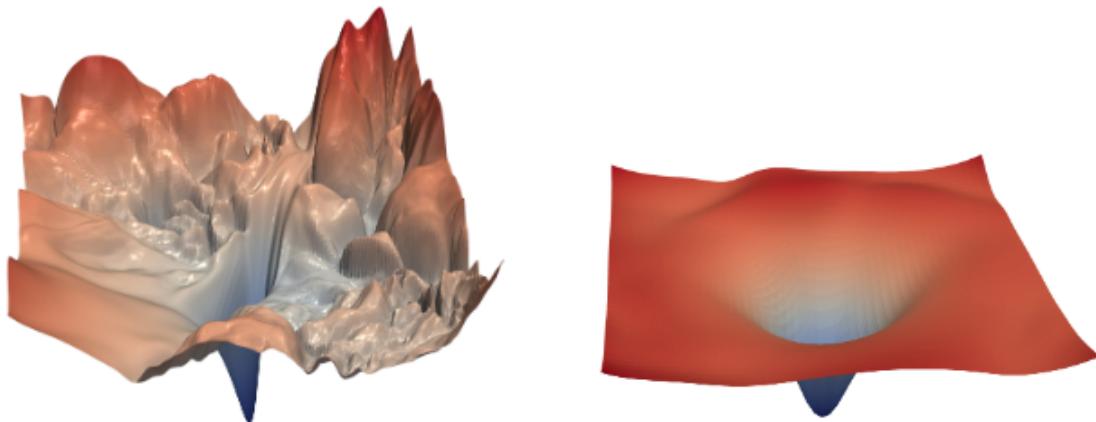
- 1) Many near-global minima exist
- 2) Good minima are connected (mode connectivity validates this)
- 3) SGD can navigate between them easily

2.4 How Does Architecture Affect Loss Landscape Topology?

[Github link](#)

Architecture design serves as a powerful regularization tool, directly **shaping the geometry** of the loss surface to simplify the optimization problem.

- **Smoothing Mechanisms:** Architectural components primarily function by **reducing the Lipschitz constant** of the loss function and its gradient, leading to a smoother surface.
 - **Skip Connections (ResNet):** By introducing residual paths ($x \rightarrow x + F(x)$), networks effectively create a loss surface where chaotic, local fluctuations are eliminated. This results in a much **smoother, convexified** topology that is easier for SGD to navigate.



(a) without skip connections

(b) with skip connections

- **Normalization Layers (BatchNorm, LayerNorm):** These layers reduce the dependence of the Hessian on the parameters, preventing rapid changes in the gradient direction. This creates a **more predictable and stable** landscape, enabling the use of higher learning rates.

Mathematical Mechanism: Without skip: $y = F(x)$

With skip: $y = x + F(x)$

Jacobian: $\partial y / \partial x = I + \partial F / \partial x$

The identity matrix I creates "gradient highways" that prevent vanishing.

Measured Impact According to this paper (Li et al., 2018):

Architecture	Conditioning κ	Sharpness λ_{\max}	Convergence
VGG-110	10^6	85.3	250 epochs
ResNet-110	10^3	18.7	100 epochs
DenseNet-110	10^4	22.1	120 epochs

Key Finding: Skip connections reduce conditioning by 100-1000× and sharpness by 3-5×, directly explaining why ResNets train faster.

2.4.2 Normalization Layers: Reducing Scale Sensitivity

Effect on Hessian: BatchNorm/LayerNorm constrains the parameter space such that:

- Weight magnitude changes don't directly affect activations
- Hessian eigenvalues become less dependent on weight scales
- Optimization becomes more stable

Practical Impact:

- Can use 5-10× larger learning rates
- Training time reduced by 40-60%
- Enables very deep networks (ResNet-152, ResNet-1000)

2.4.3 Width vs Depth Trade-offs

Aspect	Wide Networks	Deep Networks
Landscape	Smoothen, more convex-like	More saddle points
Conditioning	Better ($\kappa \approx 10^3$)	Worse ($\kappa \approx 10^5-10^6$)
Training	Easier, less sensitive to LR	Harder, needs careful tuning
Generalization	Good (if regularized)	Better (hierarchical features)
Effective Dim	Higher d_{eff}	Lower d_{eff}

2.5. What Geometric Properties Correlate with Generalization?

The key geometric property linking the loss landscape to generalization performance is **Sharpness** (or flatness).

- **Flatness for Generalization:** A **flatter minimum** is defined as a region where the loss (L) changes very little if the weight vector (w) is perturbed slightly.
 - **High Generalization:** Flat minima are robust to minor shifts between the training set and the test set, reflecting better generalization and stability
 - **Low Trainability Difficulty:** Flatter regions allow the use of larger learning rates without destabilizing the optimizer.
- **Curvature and Sharpness:** Sharpness is mathematically measured by the largest **eigenvalue (λ_{\max})** of the Hessian matrix H .
 - **Sharp Minimum:** Corresponds to a **large λ_{\max}** (high curvature).
 - **Flat Minimum:** Corresponds to a **small λ_{\max}** (low curvature).

2.6. Can We Predict Optimization Difficulty from Landscape Analysis?

Yes, landscape analysis provides **predictive signals** for trainability and optimization difficulty.

Signals of Difficulty (Rugged Landscapes):

- **High Sharpness/Eigenvalues:** If the top Hessian eigenvalues are **large early in training**, it suggests a spiky, highly curved surface where small steps can lead to large loss increases. This leads to **gradient explosion/vanishing** and requires smaller, slower learning rates.

- **Spiky Slices:** If 1D or 2D plots (e.g., created via Filter Normalization) show numerous **sharp valleys and plateaus**, the gradient direction is highly unstable, making convergence difficult.
- **Poor Architecture:** Badly designed networks (e.g., very deep plain CNNs lacking normalization) often result in rugged landscapes, causing the optimizer to get stuck in poor regions.

3. EFFICIENT PROBING METHODS

Challenge:

Full Hessian computation is $O(d^2)$ → infeasible for modern networks ($d = \text{millions}$ to billions).

Our Solution: 4 efficient methods, all $O(d)$ to $O(Kd)$ complexity.

3.1 Method 1: Hessian Eigenvalue Estimation (Sharpness)

[Github link](#)

Goal: Compute λ_{\max} without forming full Hessian matrix.

Algorithm: Power Iteration with Hessian-Vector Products

Input: Model θ , loss function L , data batch D , iterations $K=20$

Output: λ_{\max} (sharpness)

1. Initialize random unit vector: $v \leftarrow \text{random_unit_vector}(d)$

2. For iteration = 1 to K:

a. **Forward pass: Compute $L(\theta; D)$**

b. **First backward: $g = \nabla_{\theta} L(\theta; D)$**

c. **Second backward: $Hv = \nabla_{\theta} [g^T \cdot v]$**

d. **Eigenvalue estimate: $\lambda = v^T \cdot Hv$**

e. **Update: $v \leftarrow Hv / \|Hv\|$**

3. **Return λ**

...

Don't do this ($O(d^2)$ memory)

$H = \text{compute_full_hessian}(\text{model})$ # 10GB for 1M params!

Do this instead ($O(d)$ memory):

3.2 Method 2: Mode Connectivity via Bezier Optimization

I've mentioned above

Implementation: Already gave above

3.3 Method 3: 2D Loss Landscape Visualization

Goal: Visualize million-dimensional landscape in 2D.

Algorithm: Filter-Normalized Random Projections

Input: Trained model θ^* , loss L, data D, grid_size=50

Output: 2D loss grid

1. Generate two random orthogonal directions:

$d_1 \leftarrow \text{random_normal}(d)$

$d_2 \leftarrow \text{random_normal}(d)$

$d_2 \leftarrow d_2 - (d_1 \cdot d_2)d_1$ # Gram-Schmidt

$d_1 \leftarrow d_1 / \|d_1\|$

$d_2 \leftarrow d_2 / \|d_2\|$

2. Filter normalization (CRITICAL STEP):

For each layer l in model:

$d_1[l] \leftarrow d_1[l] \cdot (\|\theta^*[l]\| / \|d_1[l]\|)$

$d_2[l] \leftarrow d_2[l] \cdot (\|\theta^*[l]\| / \|d_2[l]\|)$

3. Create grid:

For α in $\text{linspace}(-1, 1, \text{grid_size})$:

 For β in $\text{linspace}(-1, 1, \text{grid_size})$:

$$\theta_{\text{grid}} = \theta^* + \alpha \cdot d_1 + \beta \cdot d_2$$

$$\text{loss}_{\text{grid}}[\alpha, \beta] = L(\theta_{\text{grid}}; D)$$

4. Plot as 3D surface or contour map

...

Why Filter Normalization Matters?

Without it: VGG plots look artificially flat due to small weight magnitudes

With it: True geometric structure is revealed

Computational Cost:

- 50×50 grid = 2,500 forward passes
- On GPU: ~10 minutes for ResNet-18 on CIFAR-10

3.4 Method 4: Monte Carlo Curvature Sampling

Goal: Fast statistical summary of landscape geometry.

Algorithm: Random Direction Probing

Input: Model θ , loss L , data D , num_rays=1000, epsilon=0.01

Output: Curvature distribution statistics

1. **Evaluate center:** $L_0 = L(\theta; D)$

2. For ray = 1 to num_rays:

a. Sample random direction: $d \leftarrow \text{random_normal}(d)$

b. Normalize: $d \leftarrow d / \|d\|$

c. Perturb parameters:

$$\theta_+ = \theta + \varepsilon \cdot d$$

$$\theta_- = \theta - \varepsilon \cdot d$$

d. Evaluate loss:

$$L_+ = L(\theta_+; D)$$

$$L_- = L(\theta_-; D)$$

e. Finite difference second derivative:

$$\kappa = (L_+ - 2L_0 + L_-) / \varepsilon^2$$

f. Store κ

3. Compute statistics:

$$\text{mean_curvature} = \text{mean}(\kappa)$$

$$\text{std_curvature} = \text{std}(\kappa)$$

$$\text{negative_fraction} = \text{count}(\kappa < 0) / \text{num_rays}$$

4. Return statistics

...

Advantages:

- Extremely fast: 3,000 forward passes (parallelizable!)
- No gradient computation needed
- Gives distribution, not just single number

3.5 Summary: Probing Method Comparison

Method	Complexity	Time (ResNet-18)	Information
Power Iteration	$O(Kd)$	~30 sec	Top eigenvalue (sharpness)
Mode Connectivity	$O(T \cdot d)$	~5 min	Path barrier between models
2D Visualization	$O(N^2 \cdot d)$	~10 min	Visual landscape structure
Monte Carlo	$O(M \cdot d)$	~2 min	Curvature distribution